

CS 5348 Operating Systems -- Project 3

The goals of this and the next projects are to let you get a better understanding on how OS works via coding and get familiar with some important Unix system calls via using them. In the following, we first discuss a program that implements a very simple OS and simulates the underlying computer, then discuss the project requirement.

1 Overview of a Simulated Simple OS (SimOS)

In SimOS, we simulate a computer system and then implement a very simple OS that manages the resources of the computer system. The computer system has components: CPU (with registers), memory, disk swap space, and clock. CPU drives memory (through load and store instructions) and other I/O devices (e.g., I/O for page fault). We simulate these by software function calls. When system starts, CPU executes the OS program, which loads in user programs upon submissions and initiates CPU registers for a user program so that CPU (sort of being arranged into) executes the user program. Some more details about the simulated computer components are discussed in 1.1, and the OS that manages the resources is discussed in 1.2. The outer layer user command processing (like shell) is discussed in 1.3.

1.1 Simulated Computer System

CPU has a set of registers (defined in `simos.h`) and it performs computation on these registers (implemented by function `cpu_execution` in `cpu.c`). During one CPU execution cycle, it fetches an instruction from memory and performs the actions according to the instruction. Besides the end-program instruction, the corresponding datum for the instruction is also fetched from memory. The instructions in our simulated computer are listed in Table 1.

Table 1. Instructions in a program.

OP Code	Parameters	System actions
2 (load)	Indx	Load from M[indx] into the AC register, indx (≥ 0 int)
3 (add)	Indx	Add current AC by M[indx]
4 (mul)	Indx	Multiply current AC by M[indx]
5 (if-goto)	indx, addr	A two-word instruction: indx is the operand of the first word; and addr is the operand of the second word. If M[indx] > 0 then goto addr, else continue to next instruction
6 (store)	Indx	Store current AC into M[indx]
7 (print)	Indx	Print the content of M[indx], AC does not change
8 (sleep)	time	Sleep for the given time in microseconds
1 (end)	Null	End of the current job, null is 0 and is unused

At the end of each instruction execution cycle, CPU checks the interrupt vector. If some interrupt bits are set, then the corresponding interrupt handling actions are performed (by `handle_interrupt` in `cpu.c`). Currently, we define 3 interrupt bits (defined in `simos.h`) and they are discussed in Table 2. The interrupt vector is initialized to 0. After completing interrupt handling, the interrupt vector is reset to 0. (What if interrupt bits get set during interrupt handling?)

Table 2. Interrupt vector.

Interrupt bit position	Description
0 for cpu time quantum (vector value = 1)	Upon activating the execution of a process, process manager sets a one-time timer of duration time quantum. When the time quantum expires, the timer sets this interrupt bit. The handler then sets the

	process to ready state and returns the control to process manager, which will switch process.
1 for age scan timer (vector value = 2)	Memory should set a periodical timer for scan and update of the memory age vectors. Timer sets this interrupt bit when the time is up. The interrupt handler simply invokes the aging scan activity for this interrupt bit.
2 for IO completion (vector value = 4)	When IO of a waiting process completes (such as page fault IO), the manager places the process into an endIO list and sets this interrupt bit. There may be multiple processes with their IO completed but the bit will only be set once. Thus, the interrupt handler fetches all processes in the endIO list and put them in the ready queue.

Memory provides CPU the functions required during CPU execution, including `get_instruction(offset)`, `get_data(offset)`, and `put_data(offset)` in `memory.c`. The parameter `offset` is based on the address space of the process and has to be converted to physical memory address (this conversion is supposed to be done above the physical memory, but for convenience, we put it in physical memory). When a new process is submitted to the system, the system will load the program and the corresponding data into memory. So memory unit also provides functions `load_instruction(...)` and `load_data(...)` for this purpose (simulating one type of direct memory access (DMA) without going through CPU).

Terminal (`term.c`) outputs a string to the monitor. When CPU (`cpu.c`) processes the print instruction, it puts what is to be printed as a string and sends the string to the terminal queue. When a process ends, process manager (`process.c`) prepares an end-program message and sends it to terminal queue. At the same time, the running process is switched into the waiting state. Since terminal is an IO device working in parallel with CPU, we need to run terminal manager as a thread. The terminal thread goes through the printing requests in the queue and process them (print the string). Note that inserting requests to the terminal queue is done by the main thread and the terminal thread removes requests from the queue. Thus, they need to be synchronized.

A data intensive process may use a lot of memory space for processing a large amount of data. It may not be possible to put all the pages of a process fully in memory. Swap space is used to allow the system to only keep currently needed pages in memory and put the remaining pages for a process on disk (called swap space). We simulate the disk swap space and implement the swap space manager in `swap.c`.

1.2 Simulated Simple OS

In `process.c`, the OS implements its process management functions. First, a PCB structure is defined in `simos.h`. A ready queue is implemented and process scheduling should be done on the ready queue. Currently the scheduling policy is a FIFO based Round Robin. When a process is submitted, the `submit_process()` function is invoked, which loads the instructions and data of the process into memory, prepares the process control block PCB, and places the process in the ready queue. Execution of processes is supposed to be done continuously under the control of OS, but we use an admin command to directly control the execution to allow easier observation of the behaviors of the system. When executing a process `execute_process()`, a process is fetched from the ready queue and the CPU function `cpu_execution` is called for executing the program.

Some code for the simulated OS is mixed with the simulated computer system. In `memory.c`, two functions `allocate_memory(...)` and `free_memory(...)` are offered, which are supposed to be implemented in the OS. In terminal, the function `terminal_output` simulates terminal printing and the other functions are terminal manager within the OS. Similarly, the swap space manager has the simulated disk IO (`read_swap_page` and `write_swap_page`) functions and the OS management code (the remaining functions) mixed in `swap.c`.

Timer provides the function (add_timer in timer.c) to let other components of the system to setup a timer event. Each timer includes the time (relative, counting from current time), the action to be performed after the time is up, and whether the timer is periodical. When the time is up, the timer manager takes the specified action. The supported actions are discussed in Table 3. Timer manager also automatically inserts the timer back if it is a periodical timer. When a timer is set, the pointer to the timer is returned (casted as an unsigned integer to avoid the necessity of exposing the internal timer structure), which can be used to locate and deactivate the timer by the deactivate_timer function. Note that in a real system, timer is supposed to have its own clock and will check the timer events upon a certain number of clock cycles. In our simulated system, we use CPU cycles as the clock.

Table 3. Actions for the timer.

Action Code	Corresponding action
1 for time quantum action	Set the corresponding interrupt bit and put the process in endWait queue in order for it to go back to the ready queue
2 for age vector scan action	Simply set the corresponding interrupt bit
3 for null action	Do nothing

1.3 Command Processing

We simulate control commands by an administrator in admin.c. Available commands to the system are listed in Table 4.

Table 4. Commands in the system.

Action	Parameters	System actions
T	-	Terminate the entire system
s (submit)	fnum	Submit a new process, should be shifted to client program
x (execute)	-	Execute a program from ready queue for one time quantum
r (register)	-	Dump registers
q (queue)	-	Dump ready queue and list of processes completed IO
p (PCB)	-	Dump PCB for every process
e (timer events)	-	Dump timer event list
m (memory)	-	Dump memory related information
d (term & swap)	-	Dump the swap request queue and the terminal queue

2 Project 3 Description

The goal for implementing the computer and OS simulator, simOS, is to let you get a deeper understanding of the activities going on in computer and OS. It may be challenging for you to implement the simOS system completely. Thus, we provide a working simulator “simos.exe” which implements simple OS resource management policies. Some components can be replaced by more advanced resource management algorithms. You will implement a more advanced memory management solution for simOS. Now, to let you get hands-on with the simulator, we removed some parts of the code and ask you to put them back in to complete the system.

cpu.c	The simulated CPU executes instructions and handles interrupts. You need to complete the missing instruction execution code (micro instructions on registers) and understand the interrupt handling functions.
process.c	The process manager manages ready queue and process execution. During process execution, there may be IO instructions, sleep instruction, page fault, or error that will stop the regular execution of the process. Also, some events may occur in the system which interrupts process execution. These are all considered in the process execution procedure. You need to put some code for the execution of an idle process to get a closer look into the process management tasks. You also need to add code for simulated context switch.

term.c	We have removed the synchronization code in this component and you need to add them back. You need to make sure that the termIO function does not busy looping on null terminal queue. You also need to achieve mutual exclusive accesses to terminal queue. Add semaphore controls to the code to achieve the goal. Make sure that your synchronization does not hurt parallelism unnecessarily.
--------	--

Files cpu.c and process.c are marked with “*** ADD CODE” strings to indicate where you have to add the desired code. In term.c, your goal is to achieve synchronization. So where to add sync code will be your design.

Another change needed to make the system work is to port simOS to your client-computer framework. In Project 1, you have Computer as a forked process interacting with Admin via pipe. In this project, we will not have Admin and Computer as two separate processes. There will only be Client and Computer interacting via sockets. You need to rewrite submit.c and term.c for the purpose. Now, simOS is the Computer. Your Computer code of Project 1 can be used in place of submit.c. But some logic in submit.c should be retained to ensure that the system will work.

You used to have two threads in Computer, one accepts new Client connections, the other reads from multiple Client sockets in turn. You can now use one thread and use “select” system call to wait for connections and read data from all clients (as shown in the sample code). Your

Your Client code in Project 1 will remain almost the same, which reads file name from the terminal and sends it to Computer. But Computer may return multiple print out messages to Client. You can use the message “Process ... terminated” as the indicator of process termination.

You also need to send the computation results to Client so that Client can print them on its own window. This requires you to change the code for function “terminal_output” in term.c, which now prints the output to a file. You need to save the socket file descriptor sockfd so that communication to the Clients can be done properly. In simOS, you can save sockfd in PCB (in real OS, PCB has a pointer pointing to an array of file descriptors). You need to add one more parameter to insert_termio function to pass sockfd. When you call insert_termio from cpu.c and process.c, the corresponding sockfd can be retrieved from PCB and passed to the terminal queue.

3 Input and Output of simOS

3.1 System Configuration

A system configuration file “config.sys” is used to set the configuration parameters of the system. The content of the configuration file is given in the following. The configuration parameters are defined in simos.h and read in system.c in the initialize_system function.

```
<max-process> <quantum> <idle-quantum>
-- max-process: max number of processes allowed
-- quantum: time quantum, given as number of instructions
-- idle-quantum: time quantum for the idle process, could be less than <quantum>
<page size> <total memory pages>
-- page size is the size of each page in physical memory
-- total memory pages is the total number of pages in the memory
<number of load pages> <max number of pages per process> <OS pages>
-- number of load pages is the #pages given to each process upon loading the process
-- max number of pages per process: used to determine the page table size for each process
-- OS size is the size of the address space of the OS which has to stay permanently in memory
-- all these sizes are in number of words
<age scan period> <terminal output time> <disk rw time>
-- age scan period specifies how frequently the system should perform aging vector scan
```

-- terminal output time and disk rw time are for simulating the slow IO (by usleep)

3.2 Program to be Executed in simOS

The programs to be submitted to simOS for execution are in a simulated machine code. The format for a program is specified as follows.

<memory size> N M

- memory size is the size of the memory the program requests, since we do not have instructions to allocate dynamic memory, the memory size for each program is fixed
- N is the number of instructions in the program
- M is the number of static data in the program

<instructions>

- following the first line, there should be N lines of instructions
- each instruction should have an opcode and an operand, both are integers in one line
- the last instruction should always be end-program, with a dummy operand
- the descriptions for opcode and operand are given in Table 1

<data>

- following the instructions are N lines of data, initialized in the program file
- every data should be initialized, for those without initial value, set it to 0 anyway

3.3 Running Original simos.exe

When executing simos.exe, you get the admin console. You can type in an admin command at the “command>” prompt, such as submitting a process, executing the processes in the system, or dump some data structure of the system.

In the original version of simos.exe, after you execute some programs, their outputs are printed in “terminal.out”. Note that the output is on terminal, so, outputs from different programs are garbled, i.e., in the order of when the print statements are issued. We specifically included the pid in the printout so that you can see clearly which outputs belong to which program. After you change “term.c”, the output should be printed on individual Client window and “terminal.out” should not be generated.

3.4 Command Line Input for your New Code

Your original client code has command line input Client-ID, and server host name and server port. You will not need Client-ID for simOS, since it is replaced by Pid given by the Computer program. So, your client code should be started like:

Client.exe Server-host-name Server-port-number

Your Admin.exe from Project 1 will now be simos.exe. To start your new simos.exe, you still need to supply the port-number (follow the same convention to avoid conflicts), but you will not need to have a sleep time. The command line input for simos.exe is as follows:

simos.exe server-port-number

When the Client submits the program to the Computer (simOS), the message should only contain a file name, no Client-ID. The printout sent from Computer to Client will be the same as how it is now, including the Pid and whatever to be printed. Client just printout whatever it receives without needing to add more information.

4 Submission

You need to electronically submit your program **before** midnight of the due date (check the web page for the due date). Your submission should be a zip file or tar file including the following

- ✓ All source code files that have been altered, at least including `cpu.c`, `process.c`, `term.c`, and `submit.c`.
- ✓ Source files for socket communication between Computer and Clients.
- ✓ The makefile to generate executable “simos.exe”.
- ✓ The *DesignDoc* file that contains the description of the major features of your program that is not specified in the project specification, including
 - ✓ The list of program files you have altered, and the changes you made in each program file that are beyond the “**** ADD CODE” segments,
 - ✓ How synchronization is done in your “term.c”,
 - ✓ How you changed Project 1 code for Project 3, including whether you used your own code or the posted code and the changes you made,
 - ✓ Any other issues related to the program you are submitting.
- ✓ You should submit your project through elearning.